TD3 for Competitive Air Hockey

RL Course WS 2025/26 — Final Project

Julian Jurcevic · Team: alphabet-td3

February 2026

Air Hockey Environment

Environment

- 2-player Air Hockey (Gymnasium + Box2D)
- 4 continuous actions: x, y, rotation, shoot
- 18-dim state: positions, velocities, angles
- Reward: ± 10 + shaped proximity reward

Key Challenges

- Non-stationarity: opponent improves
- Fast dynamics + sparse rewards
- Overfitting to single opponent
- Must combine attack & defense

TD3 — Three Key Ideas

1. Clipped Double Q-Learning

Two critics; target uses the *minimum*:

$$y = r + \gamma(1-d) \min_{i=1,2} Q_{\phi_i'}(s', \tilde{a}')$$

2. Target Policy Smoothing

$$\tilde{a}' = \text{clip}\big(\pi_{\theta'}(s') + \epsilon,\ -c,\ c\big)$$

3. Delayed Actor Updates

- Critic: every step; Actor: every 2nd
- Polyak averaging:
  $\phi' \leftarrow \tau\,\phi + (1-\tau)\,\phi'$

Architecture

- Actor & twin critics: 2×256 (tanh)
- Actor: $a \in [-1, 1]^4$
- Critics: concatenated (s, a) input

Exploration Noise & Replay

Noise Annealing

$$\sigma_t = \max\left(\sigma_0\left(1 - \frac{t}{T}\right),\ \sigma_{\min}\right)$$

- Broad exploration early, exploitation later
- Compared: Gaussian, OU, Pink, Uniform
- **Best: OU noise** (temporal correlation)

Replay & Hyperparameters

- Uniform replay (300k buffer)
- PER tested but hurt performance

γ	0.99
LR (actor/critic)	2×10^{-4}
τ	0.005
Batch size	256
Target noise / clip	0.2 / 0.3

Three-Stage Curriculum with Self-Play

Self-Play Pool: Snapshots every k=150 eps · Pool N_{pool}=25 · Difficulty-weighted sampling ($\times$1.2 loss, $\times$0.95 win)

Ablation: Noise Comparison (3 seeds, Stage II)

Noise Type	WR Weak (%)	WR Strong (%)	Ret. Weak	Ret. Strong
Gaussian	92.5 ± 4.5	81.0 ± 0.5	8.22 ± 0.80	5.69 ± 0.10
Ornstein–Uhlenbeck	$\mathbf{94.7 \pm 0.6}$	$\mathbf{89.0 \pm 2.7}$	$\mathbf{8.56 \pm 0.13}$	$\mathbf{7.06 \pm 0.46}$
Pink	92.6 ± 4.3	86.1 ± 2.8	8.17 ± 0.57	6.40 ± 0.34
Uniform	91.2 ± 2.8	80.2 ± 8.4	8.10 ± 0.55	5.51 ± 1.51

- OU best: **+8%** vs. Gaussian against strong (temporal correlation $\rightarrow$ smoother trajectories)
- Pink noise second-best (also correlated) · Uniform: highest variance

Ablation: Self-Play & Prioritized Replay (3 seeds)

Variant	WR Weak (%)	WR Strong (%)	Ret. Weak	Ret. Strong
No PER, No SP	**93.1 ± 3.8**	**78.3 ± 3.1**	**8.33 ± 0.66**	**5.00 ± 0.70**
No PER, Self-Play	90.7 ± 5.9	72.6 ± 7.6	7.62 ± 1.26	4.06 ± 1.56
PER, No SP	75.8 ± 9.2	66.1 ± 4.7	4.22 ± 2.00	1.99 ± 1.04
PER, Self-Play	78.3 ± 2.2	65.3 ± 5.1	4.71 ± 0.54	1.78 ± 1.01

PER hurts performance

- Non-stationary opponents amplify priority variance
- ≈15–20% WR drop ⇒ **not used**

Self-Play — trade-off

- Lower benchmark scores but **retained** for tournament robustness
- Pool diversity decreases as policies converge

Curriculum Training Progression

Training curves for a representative single-seed run (see report Figure 1).

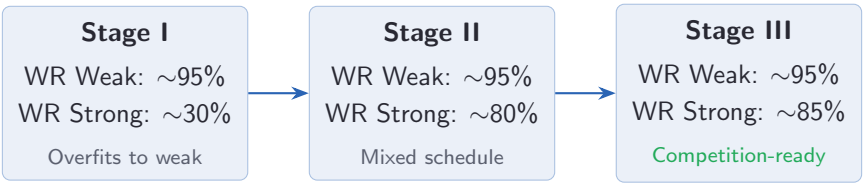

- Without curriculum: strong win rate stays at ∼30% (pure weak training)
- Staged scheduling resolves this while retaining high weak win rates
- Model selection via $\min(\mathrm{WR}_{\mathrm{weak}}, \mathrm{WR}_{\mathrm{strong}})$ enforces robustness

Conclusion & Takeaways

What worked

- **Curriculum**: most impactful; prevents overfitting to single opponent
- **OU noise**: +8% WR (strong) via temporal correlation
- **Self-play**: retained for tournament generalization

What didn't work

- **PER**: ∼15% drop under non-stationarity

Final Agent Config

Algorithm	TD3
Noise	OU (annealed)
Replay	Uniform (300k)
Curriculum	3 stages
Self-Play	Pool of 25
WR Weak	∼95%
WR Strong	∼85%

Limitations

- Single-seed training curves
- Manual curriculum tuning
- Self-play pool convergence

Thank you!

Questions?